



# Data Science

## Association Rules

Themis Palpanas  
University of Paris

Data Science

1

1

### Thanks for slides to:



- Jiawei Han
- Jeff Ullman

Data Science

2

2

# Roadmap

---

- Frequent Patterns
  - Frequent Pattern Analysis
  - Applications
    - Market-Basket Model
    - Association Rules
- A-Priori Algorithm
- Improvements to A-Priori

Data Science

3

3

## What Is Frequent Pattern Analysis?

---

- **Frequent pattern**: a pattern (a set of items, subsequences, substructures, etc.) that occurs frequently in a data set
- First proposed by Agrawal, Imielinski, and Swami [AIS93] in the context of **frequent itemsets** and **association rule mining**
- Motivation: Finding inherent regularities in data
  - What products were often purchased together?— Beer and diapers?!
  - What are the subsequent purchases after buying a PC?
  - What kinds of DNA are sensitive to this new drug?
  - Can we automatically classify web documents?
- Applications
  - Basket data analysis, cross-marketing, catalog design, sale campaign analysis, Web log (click stream) analysis, and DNA sequence analysis.

Data Science

4

4

## Why Is Freq. Pattern Mining Important?

- Discloses an intrinsic and important property of data sets
- Forms the foundation for many essential data mining tasks
  - Association, correlation, and causality analysis
  - Sequential, structural (e.g., sub-graph) patterns
  - Pattern analysis in spatiotemporal, multimedia, time-series, and stream data
  - Classification: associative classification
  - Cluster analysis: frequent pattern-based clustering
  - Data warehousing: iceberg cube and cube-gradient
  - Semantic data compression: fascicles
  - Broad applications

Data Science

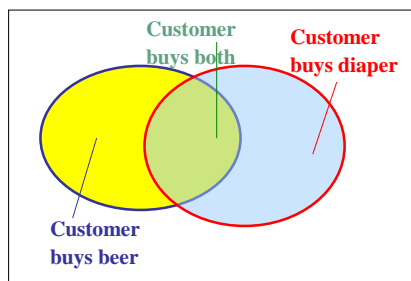
5

5

## Basic Concepts: Frequent Patterns and Association Rules

Transaction-id	Items bought
10	A, B, D
20	A, C, D
30	A, D, E
40	B, E, F
50	B, C, D, E, F

- Itemset  $X = \{x_1, \dots, x_k\}$
- Find all the rules  $X \rightarrow Y$  with minimum support and confidence
  - **support**,  $s$ , **probability** that a transaction contains  $X \cup Y$
  - **confidence**,  $c$ , **conditional probability** that a transaction having  $X$  also contains  $Y$



Let  $sup_{min} = 50\%$ ,  $conf_{min} = 50\%$   
 Freq. Pat.:  $\{A:3, B:3, D:4, E:3, AD:3\}$   
 Association rules:  
 $A \rightarrow D$  (60%, 100%)  
 $D \rightarrow A$  (60%, 75%)

Data Science

6

6

# The Market-Basket Model

---

- A large set of *items*, e.g., things sold in a supermarket.
- A large set of *baskets*, each of which is a small set of the items, e.g., the things one customer buys on one day.

Data Science

7

7

## Support

---

- Simplest question: find sets of items that appear “frequently” in the baskets.
- *Support* for itemset  $I$  = the number of baskets containing all items in  $I$ .
- Given a support *threshold*  $s$ , sets of items that appear in  $\geq s$  baskets are called *frequent itemsets*.

Data Science

8

8

## Example

- Items={milk, coke, pepsi, beer, juice}.
- Support = 3 baskets.

$B_1 = \{m, c, b\}$	$B_2 = \{m, p, j\}$
$B_3 = \{m, b\}$	$B_4 = \{c, j\}$
$B_5 = \{m, p, b\}$	$B_6 = \{m, c, b, j\}$
$B_7 = \{c, b, j\}$	$B_8 = \{b, c\}$
- Frequent itemsets?

Data Science

9

9

## Example

- Items={milk, coke, pepsi, beer, juice}.
- Support = 3 baskets.

$B_1 = \{m, c, b\}$	$B_2 = \{m, p, j\}$
$B_3 = \{m, b\}$	$B_4 = \{c, j\}$
$B_5 = \{m, p, b\}$	$B_6 = \{m, c, b, j\}$
$B_7 = \{c, b, j\}$	$B_8 = \{b, c\}$
- Frequent itemsets:
  - $\{m\}, \{c\}, \{b\}, \{j\}, \{m, b\}, \{c, b\}, \{j, c\}$ .

Data Science

10

10

## Applications --- (1)

---

- **Real market baskets:** chain stores keep terabytes of information about what customers buy together.
  - Tells how typical customers navigate stores, lets them position tempting items.
  - Suggests tie-in “tricks,” e.g., run sale on diapers and raise the price of beer.
- High support needed, or no \$\$’s .

Data Science

11

11

## Applications --- (2)

---

- “Baskets” = documents; “items” = words in those documents.
  - Lets us find words that appear together unusually frequently, i.e., linked concepts.
- “Baskets” = sentences, “items” = documents containing those sentences.
  - Items that appear together too often could represent plagiarism.

Data Science

12

12

## Applications --- (3)

---

- “Baskets” = Web pages; “items” = linked pages.
  - Pairs of pages with many common references may be about the same topic.
- “Baskets” = Web pages  $p$ ; “items” = pages that link to  $p$ .
  - Pages with many of the same links may be mirrors or about the same topic.

Data Science

13

13

## Important Point

---

- “Market Baskets” is an abstraction that models any many-many relationship between two concepts: “items” and “baskets.”
  - Items need not be “contained” in baskets.
- The only difference is that we count co-occurrences of items related to a basket, not vice-versa.

Data Science

14

14

## Scale of Problem

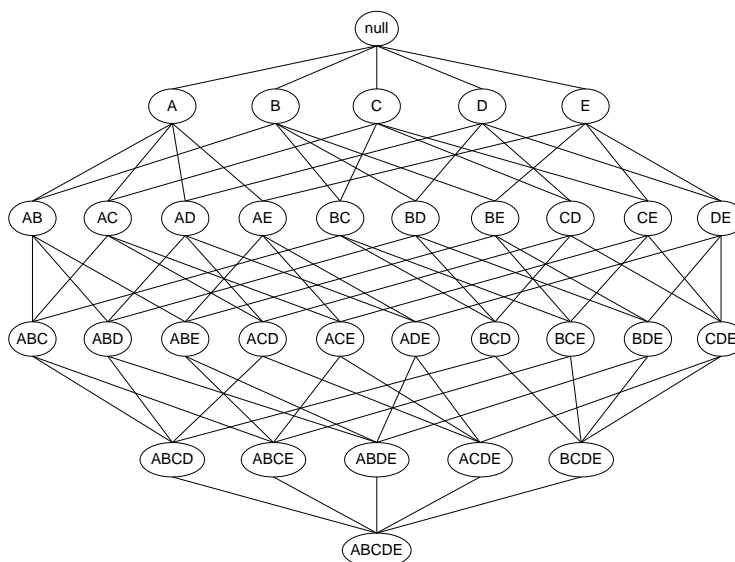
- WalMart sells 100,000 items and can store billions of baskets.
- The Web has over 100,000,000 words and billions of pages.

Data Science

15

15

## A simple algorithm for finding all frequent itemsets ??



Data Science

16

16



## Association Rules

- If-then rules about the contents of baskets.
- $\{i_1, i_2, \dots, i_k\} \rightarrow j$  means: "if a basket contains all of  $i_1, \dots, i_k$  then it is *likely* to contain  $j$ ."
- *Confidence* of this association rule is the probability of  $j$  given  $i_1, \dots, i_k$

Data Science

17

17

## Example

$$\begin{array}{ll} B_1 = \{m, c, b\} & B_2 = \{m, p, j\} \\ B_3 = \{m, b\} & B_4 = \{c, j\} \\ B_5 = \{m, p, b\} & B_6 = \{m, c, b, j\} \\ B_7 = \{c, b, j\} & B_8 = \{b, c\} \end{array}$$

- An association rule:  $\{m, b\} \rightarrow c$ .
  - Confidence = ?

Data Science

18

18

## Example

$B_1 = \{m, c, b\}$        $B_2 = \{m, p, j\}$   
 $B_3 = \{m, b\}$        $B_4 = \{c, j\}$   
 $B_5 = \{m, p, b\}$        $B_6 = \{m, c, b, j\}$   
 $B_7 = \{c, b, j\}$        $B_8 = \{b, c\}$

- An association rule:  $\{m, b\} \rightarrow c$ .
  - Confidence =  $2/4 = 50\%$ .

Data Science

19

19

## Interest

- The *interest* of an association rule  $X \rightarrow Y$  is the absolute value of the amount by which the confidence differs from the probability of  $Y$ .

Data Science

20

20

## Example

$$\begin{array}{ll} B_1 = \{m, c, b\} & B_2 = \{m, p, j\} \\ B_3 = \{m, b\} & B_4 = \{c, j\} \\ B_5 = \{m, p, b\} & B_6 = \{m, c, b, j\} \\ B_7 = \{c, b, j\} & B_8 = \{b, c\} \end{array}$$

- For association rule  $\{m, b\} \rightarrow c$ , item  $c$  appears in 5/8 of the baskets.
- Interest =  $|2/4 - 5/8| = 1/8$  --- not very interesting.

Data Science

21

21

## Relationships Among Measures

- Rules with high support and confidence may be useful even if they are not “interesting.”
  - We don’t care if buying bread *causes* people to buy milk, or whether simply a lot of people buy both bread and milk.
- But high interest suggests a cause that might be worth investigating.

Data Science

22

22

## Finding Association Rules

---

- A typical question: “find all association rules with support  $\geq s$  and confidence  $\geq c$ .”
  - Note: “support” of an association rule is the support of the set of items it mentions.
- Hard part: finding the high-support (*frequent*) itemsets.
  - Checking the confidence of association rules involving those sets is relatively easy.

Data Science

23

23

## Computation Model

---


- Typically, data is kept in a “flat file” rather than a database system.
  - Stored on disk.
  - Stored basket-by-basket.
  - Expand baskets into pairs, triples, etc. as you read baskets.

Data Science

24

24

## File Organization



Data Science

25

25

## Computation Model --- (2)

- The true cost of mining disk-resident data is usually the **number of disk I/O's**.
- In practice, association-rule algorithms read the data in **passes** --- all baskets read in turn.
- Thus, we measure the cost by the number of passes an algorithm takes.

Data Science

26

26

## Main-Memory Bottleneck

---

- For many frequent-itemset algorithms, main memory is the critical resource.
  - As we read baskets, we need to count something, e.g., occurrences of pairs.
  - The number of different things we can count is limited by main memory.
  - Swapping counts in/out is a disaster.

Data Science

27

27

## Finding Frequent Pairs

---

- The hardest problem often turns out to be finding the frequent pairs.
- We'll concentrate on how to do that, then discuss extensions to finding frequent triples, etc.

Data Science

28

28

## Naïve Algorithm

- Read file once, counting in main memory the occurrences of each pair.
  - Expand each basket of  $n$  items into its  $n(n-1)/2$  pairs.
- Fails if  $(\text{\#items})^2$  exceeds main memory.
  - **Remember:** #items can be 100K (Wal-Mart) or 10B (Web pages).

Data Science

29

29

## Details of Main-Memory Counting


- **Two approaches:**
  1. Count all item pairs, using a triangular matrix.
  2. Keep a table of triples  $[i, j, c]$  = the count of the pair of items  $\{i, j\}$  is  $c$ .
- (1) requires only (say) 4 bytes/pair.
- (2) requires 12 bytes, but only for those pairs with count  $> 0$ .

Data Science

30

30

## Details of Main-Memory Counting



Data Science

31

31

## Details of Approach #1

- Number items 1, 2,...
- Keep pairs in the order  $\{1,2\}, \{1,3\}, \dots, \{1,n\}, \{2,3\}, \{2,4\}, \dots, \{2,n\}, \{3,4\}, \dots, \{3,n\}, \dots, \{n-1,n\}$ .
- Find pair  $\{i, j\}$  at the position:
  - $(i-1)(n-i/2) + j - i$
- Total number of pairs  $n(n-1)/2$ ; total bytes about  $2n^2$ .

Data Science

32

32



## Details of Approach #2

---

- You need a hash table, with  $i$  and  $j$  as the key, to locate  $(i, j, c)$  triples efficiently.
  - Typically, the cost of the hash structure can be neglected.
- Total bytes used is about  $12p$ , where  $p$  is the number of pairs that actually occur.
  - Beats triangular matrix if at most  $1/3$  of possible pairs actually occur.

Data Science

33

33

## Roadmap

---

- Frequent Patterns
- A-Priori Algorithm
  - Monotonicity Property
  - Algorithm Description
- Improvements to A-Priori

Data Science

34

34

## A-Priori Algorithm --- (1)

- A two-pass approach called *a-priori* limits the need for main memory.
- Key idea: *monotonicity*: if a set of items appears at least  $s$  times, so does every subset.
  - *Contrapositive for pairs*: if item  $i$  does not appear in  $s$  baskets, then no pair including  $i$  can appear in  $s$  baskets.

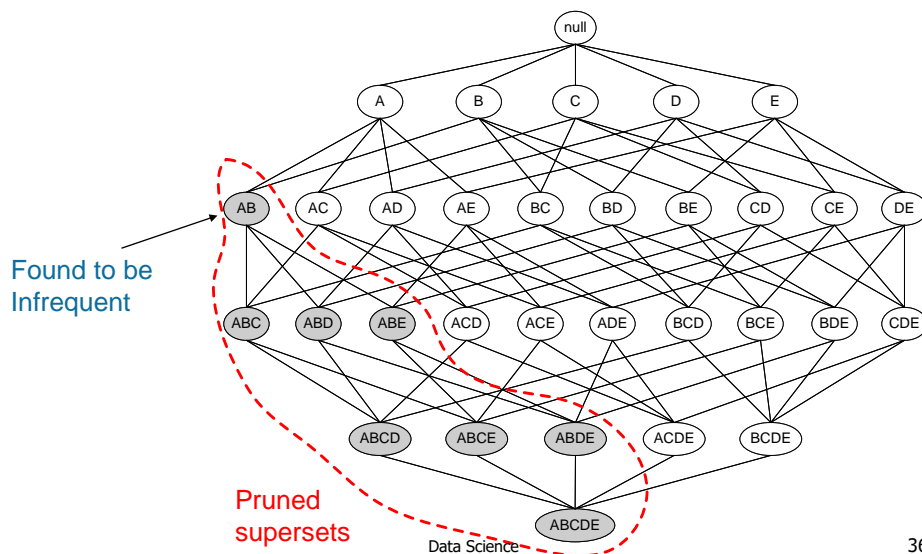
(Agrawal & Srikant @VLDB'94, Mannila, et al. @ KDD' 94)

Data Science

35

35

## Illustrating the Apriori principle



Data Science

36

36

## A-Priori Algorithm --- (1)

---



Data Science

37

37

## A-Priori Algorithm --- (1)

---



Data Science

38

38

## A-Priori Algorithm --- (1)

---



Data Science

39

39

## A-Priori Algorithm --- (1)

---



Data Science

40

40

## A-Priori Algorithm --- (2)

---

- **Pass 1:** Read baskets and count in main memory the occurrences of each item.
  - memory requirements?

Data Science

41

41

## A-Priori Algorithm --- (2)

---

- **Pass 1:** Read baskets and count in main memory the occurrences of each item.
  - Requires only memory proportional to #items.

Data Science

42

42

## A-Priori Algorithm --- (2)

---

- **Pass 1:** Read baskets and count in main memory the occurrences of each item.
  - Requires only memory proportional to #items.
- **Pass 2:** Read baskets again and count in main memory only those pairs both of which were found in Pass 1 to be frequent.
  - memory requirements?

Data Science

43

43

## A-Priori Algorithm --- (2)

---

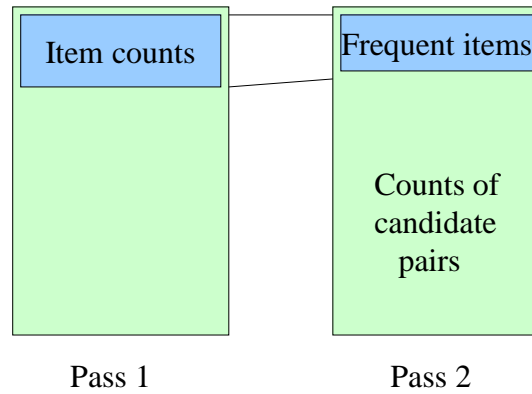
- **Pass 1:** Read baskets and count in main memory the occurrences of each item.
  - Requires only memory proportional to #items.
- **Pass 2:** Read baskets again and count in main memory only those pairs both of which were found in Pass 1 to be frequent.
  - Requires memory proportional to square of frequent items only.

Data Science

44

44

## Picture of A-Priori

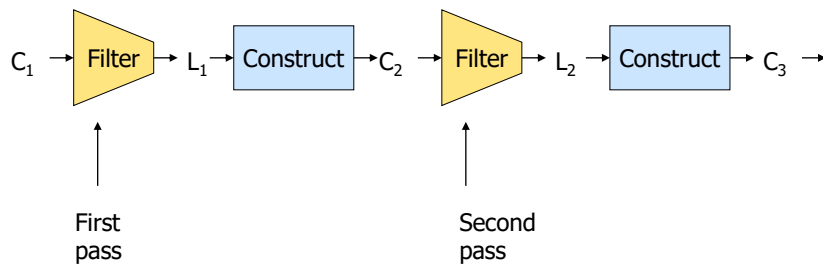


Data Science

45

45

## Frequent Triples, Etc.



Data Science

46

46

## A-Priori for All Frequent Itemsets

- One pass for each  $k$ .
- Needs room in main memory to count each candidate  $k$ -tuple.
- For typical market-basket data and reasonable support (e.g., 1%),  $k = 2$  requires the most memory.

Data Science

47

47

## The Apriori Algorithm—An Example

Database TDB  $\text{Sup}_{\min} = 2$

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

Data Science

48

48



## The Apriori Algorithm—An Example

Database TDB

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

$\text{Sup}_{\min} = 2$

$C_1$

1<sup>st</sup> scan

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

Data Science

49

49

## The Apriori Algorithm—An Example

Database TDB

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

$\text{Sup}_{\min} = 2$

$C_1$

1<sup>st</sup> scan

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

$L_1$

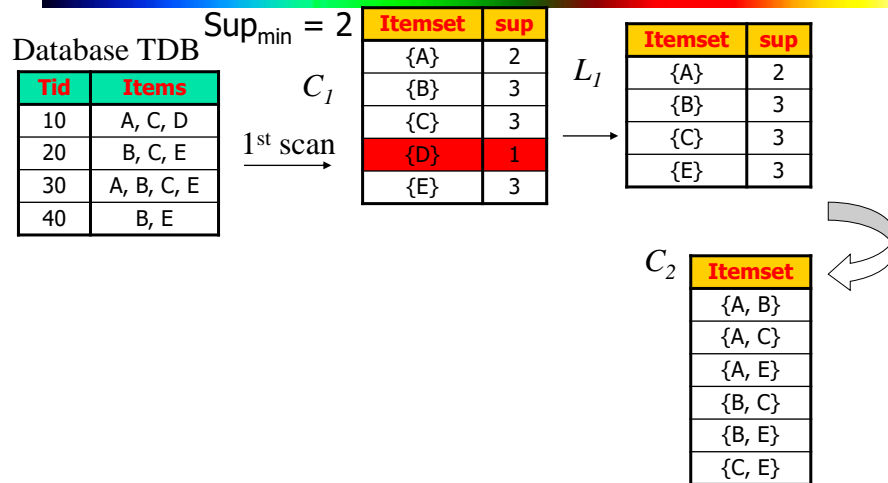
Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3

Data Science

50

50

## The Apriori Algorithm—An Example

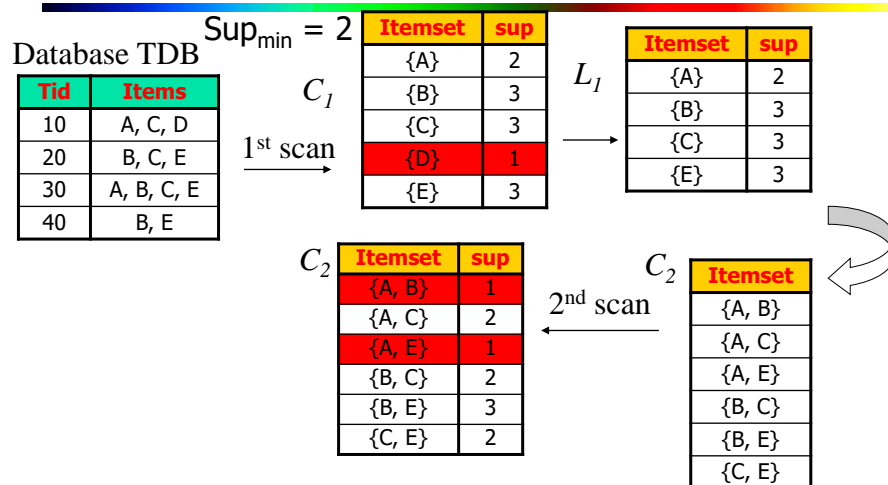


Data Science

51

51

## The Apriori Algorithm—An Example

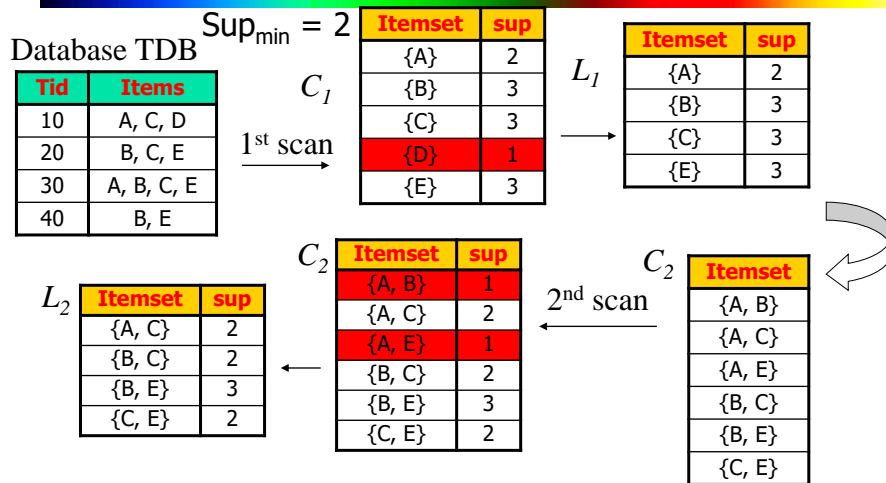


Data Science

52

52

## The Apriori Algorithm—An Example

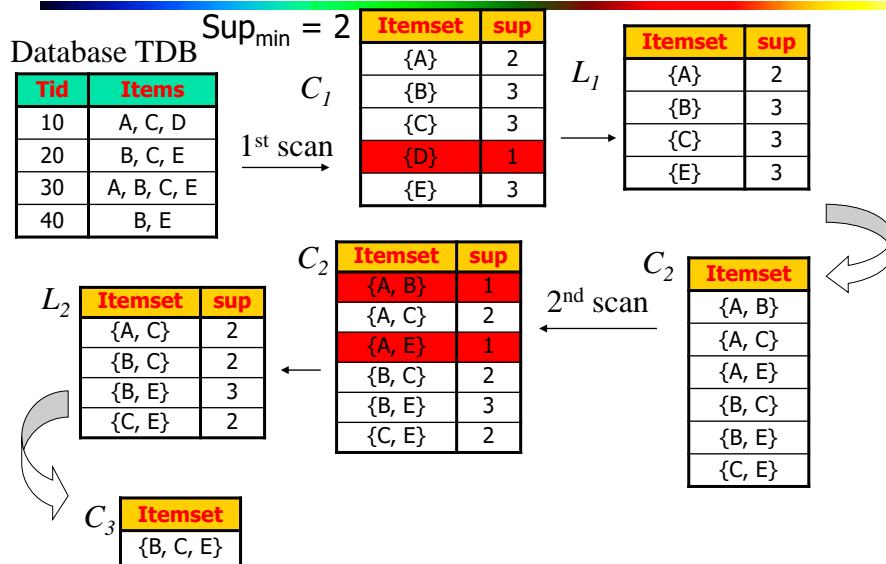


Data Science

53

53

## The Apriori Algorithm—An Example

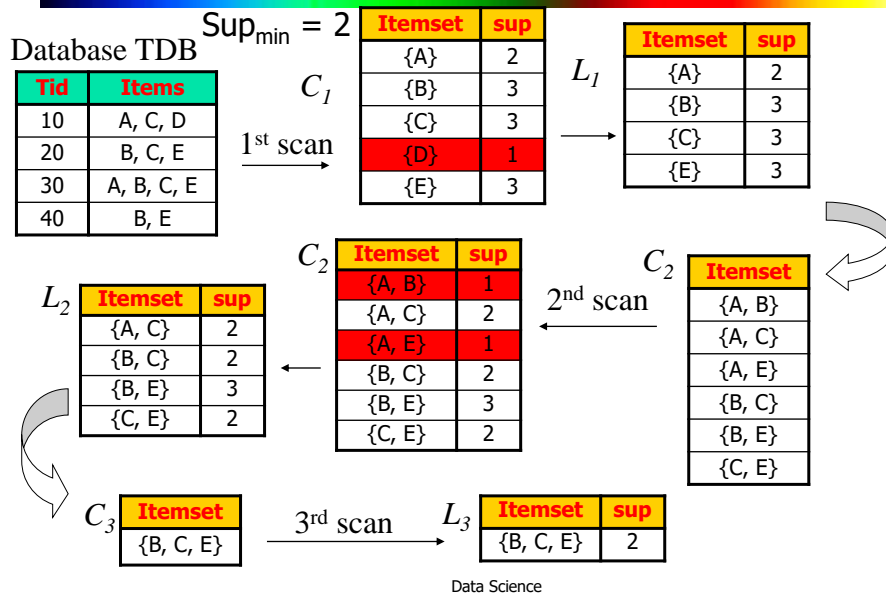


Data Science

54

54

## The Apriori Algorithm—An Example



55

55

## The Apriori Algorithm

- Pseudo-code:

$C_k$ : Candidate itemset of size k

$L_k$ : frequent itemset of size k

$L_1 = \{\text{frequent items}\};$

**for** ( $k = 1; L_k \neq \emptyset; k++$ ) **do begin**

$C_{k+1}$  = candidates generated from  $L_k$ ;

**for each** transaction  $t$  in database **do**

increment the count of all candidates in  $C_{k+1}$   
that are contained in  $t$

$L_{k+1}$  = candidates in  $C_{k+1}$  with min\_support

**end**

**return**  $\cup_k L_k$ ;

Data Science

56

56

## Important Details of Apriori

- How to generate candidates?
  - Step 1: self-joining  $L_k$
  - Step 2: pruning
- How to count supports of candidates?
- Example of Candidate-generation
  - $L_3 = \{abc, abd, acd, ace, bcd\}$
  - Self-joining:  $L_3 * L_3$ 
    - $abcd$  from  $abc$  and  $abd$
    - $acde$  from  $acd$  and  $ace$
  - Pruning:
    - $acde$  is removed because  $ade$  is not in  $L_3$
  - $C_4 = \{abcd\}$

Data Science

57

57

## How to Generate Candidates?

- Suppose the items in  $L_{k-1}$  are listed in an order
- Step 1: self-joining  $L_{k-1}$ 
  - insert into  $C_k$
  - select  $p.item_1, p.item_2, \dots, p.item_{k-1}, q.item_{k-1}$
  - from  $L_{k-1} p, L_{k-1} q$
  - where  $p.item_1 = q.item_1, \dots, p.item_{k-2} = q.item_{k-2}, p.item_{k-1} < q.item_{k-1}$
- Step 2: pruning
  - for all **itemsets**  $c$  in  $C_k$  do
  - for all **(k-1)-subsets**  $s$  of  $c$  do
  - if ( $s$  is not in  $L_{k-1}$ ) then delete  $c$  from  $C_k$

Data Science

58

58

## How to Count Supports of Candidates?

- Why counting supports of candidates a problem?
  - The total number of candidates can be huge
  - One transaction may contain many candidates
- Method:
  - Candidate itemsets are stored in a *hash-tree*
  - *Leaf node* of hash-tree contains a list of itemsets and counts
  - *Interior node* contains a hash table
  - *Subset function*: finds all the candidates contained in a transaction

Data Science

59

59

## Exploiting the Apriori principle

1. Find **frequent 1-items** and put them to  $L_k$  ( $k=1$ )
2. Use  $L_k$  to generate a collection of *candidate* itemsets  $C_{k+1}$  with size ( $k+1$ )
3. Scan the database to find which itemsets in  $C_{k+1}$  are **frequent** and put them into  $L_{k+1}$
4. If  $L_{k+1}$  is not empty
  - $k=k+1$
  - Goto step 2

R. Agrawal, R. Srikant: "Fast Algorithms for Mining Association Rules",  
*Proc. of the 20th Int'l Conference on Very Large Databases*, 1994.

60

60

# The Apriori algorithm

$C_k$ : Candidate itemsets of size k

$L_k$ : frequent itemsets of size k

$L_1 = \{\text{frequent 1-itemsets}\};$

for ( $k = 2; L_k \neq \emptyset; k++$ )

$C_{k+1} = \text{GenerateCandidates}(L_k)$

for each transaction  $t$  in database do

increment count of candidates in  $C_{k+1}$  that are contained in  $t$

endfor

$L_{k+1} = \text{candidates in } C_{k+1} \text{ with support } \geq \text{min\_sup}$

endfor

return  $\bigcup_k L_k$

61

61

## GenerateCandidates

- Assume the items in  $L_k$  are listed in an order (e.g., alphabetical)
- Step 1: self-joining  $L_k$  (IN SQL)**

insert into  $C_{k+1}$

select  $p.item_1, p.item_2, \dots, p.item_k, q.item_k$

from  $L_k p, L_k q$

where  $p.item_1 = q.item_1, \dots, p.item_{k-1} = q.item_{k-1}, p.item_k < q.item_k$


62

62

## Example of Candidates Generation

- $L_3 = \{abc, abd, acd, ace, bcd\}$
- **Self-joining:**  $L_3 * L_3$

- $abcd$  from  $abc$  and  $abd$
- $acde$  from  $acd$  and  $ace$



63

63

## GenerateCandidates

- Assume the items in  $L_k$  are listed in an order (e.g., alphabetical)
- **Step 1: self-joining  $L_k$  (IN SQL)**
  - insert into  $C_{k+1}$
  - select  $p.item_1, p.item_2, \dots, p.item_k, q.item_k$
  - from  $L_k p, L_k q$
  - where  $p.item_1 = q.item_1, \dots, p.item_{k-1} = q.item_{k-1}, p.item_k < q.item_k$
- **Step 2: pruning**
  - forall **itemsets**  $c$  in  $C_{k+1}$  do
  - forall **k-subsets**  $s$  of  $c$  do
  - if ( $s$  is not in  $L_k$ ) then delete  $c$  from  $C_{k+1}$

64

64



## Example of Candidates Generation

- $L_3 = \{abc, abd, acd, ace, bcd\}$


- **Self-joining:**  $L_3 * L_3$

- $abcd$  from  $abc$  and  $abd$
- $acde$  from  $acd$  and  $ace$

- **Pruning:**

- $acde$  is removed because  $ade$  is not in  $L_3$

- $C_4 = \{abcd\}$



65

65

## The Apriori algorithm

$C_k$ : Candidate itemsets of size  $k$

$L_k$ : frequent itemsets of size  $k$

$L_1 = \{\text{frequent items}\};$

**for** ( $k = 1; L_k \neq \emptyset; k++$ )

$C_{k+1} = \text{GenerateCandidates}(L_k)$

**for** each transaction  $t$  in database **do**

increment count of candidates in  $C_{k+1}$  that are contained in  $t$

**endfor**

$L_{k+1} = \text{candidates in } C_{k+1} \text{ with support } \geq \text{min\_sup}$

**endfor**

**return**  $\bigcup_k L_k$

66

66

# How to Count Supports of Candidates?

- Naive algorithm?

## – Method:


- Candidate itemsets are stored in a *hash-tree*
- *Leaf node* of hash-tree contains a list of itemsets and counts
- *Interior node* contains a hash table
- *Subset function*: finds all the candidates contained in a transaction

67

67

## Example of the hash-tree for $C_3$

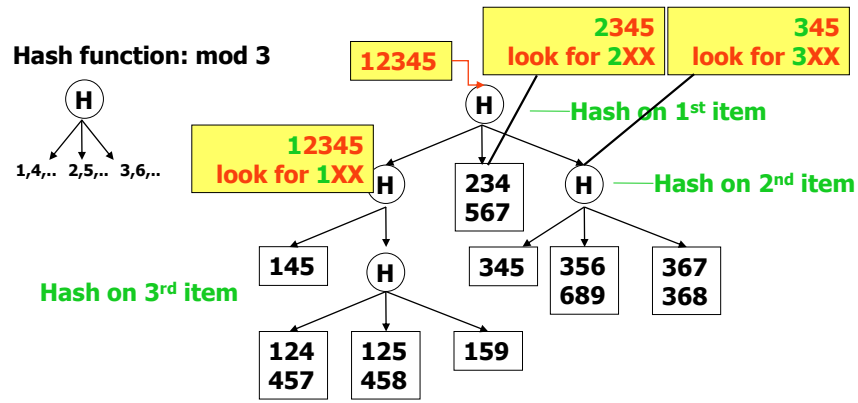
Hash function: mod 3



68

68

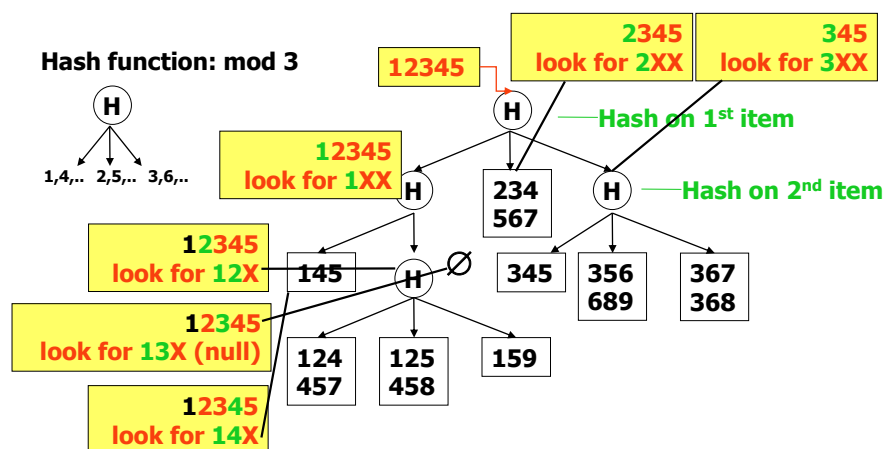
## Example of the hash-tree for $C_3$



69

69

## Example of the hash-tree for $C_3$



The subset function finds all the candidates contained in a transaction:

- At the root level it hashes on all items in the transaction
- At level  $i$  it hashes on all items in the transaction that come after item the  $i$ -th item

70

## Where are the Association Rules?

---

- so far we have seen how A-priori efficiently computes all the frequent itemsets
- but how are the association rules generated from the frequent itemsets?

Data Science

71

71

## Association Rule Generation

---

- given the frequent itemsets, generate association rules as follows
  - for each frequent itemset  $I$ 
    - generate all non-empty subsets of  $I$
  - for each non-empty subset  $s$  of  $I$ 
    - output association rule:  $s \rightarrow (I-s)$

Data Science

72

72

## Association Rule Generation

---

- given the frequent itemsets, generate association rules as follows
  - for each frequent itemset  $I$ 
    - generate all non-empty subsets of  $I$
  - for each non-empty subset  $s$  of  $I$ 
    - output association rule:  $s \rightarrow (I-s)$ , if  $\text{supp}(I)/\text{supp}(s) \geq c$

Data Science

73

73

## Association Rule Generation

---

- given the frequent itemsets, generate association rules as follows
  - for each frequent itemset  $I$ 
    - generate all non-empty subsets of  $I$
  - for each non-empty subset  $s$  of  $I$ 
    - output association rule:  $s \rightarrow (I-s)$ , if  $\text{supp}(I)/\text{supp}(s) \geq c$
- we know  $\text{supp}(\text{rule}) \geq s$ 
  - generated from frequent itemsets

Data Science

74

74

# Roadmap



- Frequent Patterns
- A-Priori Algorithm
- Improvements to A-Priori
  - Park-Chen-Yu Algorithm
  - Multistage Algorithm
  - Approximate Algorithms
  - Compacting Results

Data Science

75

75

# PCY Algorithm



- Hash-based improvement to A-Priori.
- During Pass 1 of A-priori, most memory is idle.
- Use that memory to keep counts of buckets into which pairs of items are hashed.
  - Just the count, not the pairs themselves.
- Gives extra condition that candidate pairs must satisfy on Pass 2.
- J. Park, M. Chen, and P. Yu. [An effective hash-based algorithm for mining association rules](#). In *SIGMOD'95*

Data Science

76

76

## PCY Algorithm --- Before Pass 1 Organize Main Memory

---

- Space to count each item.
  - One (typically) 4-byte integer per item.
- Use the rest of the space for as many integers, representing buckets, as we can.

Data Science

77

77

## PCY Algorithm --- Pass 1

---

```
FOR (each basket) {  
  FOR (each item)  
    add 1 to item's count;  
  FOR (each pair of items) {  
    hash the pair to a bucket;  
    add 1 to the count for that  
    bucket  
  }  
}
```

Data Science

78

78

## Observations About Buckets

1. If a bucket contains a frequent pair, then the bucket is surely frequent.
  - We cannot use the hash table to eliminate any member of this bucket.
2. Even without any frequent pair, a bucket can be frequent.
  - Again, nothing in the bucket can be eliminated.
3. But in the best case, the count for a bucket is less than the support  $s$ .
  - Now, all pairs that hash to this bucket can be eliminated as candidates, even if the pair consists of two frequent items.

Data Science

79

79

## PCY Algorithm --- Between Passes

- Replace the buckets by a bit-vector:
  - 1 means the bucket count exceeds the support  $s$  (**frequent bucket**); 0 means it did not.
- Integers are replaced by bits, so the bit-vector requires little second-pass space.
- Also, decide which items are frequent and list them for the second pass.


Data Science

80

80



## Picture of PCY



Data Science

81

81

## PCY Algorithm --- Pass 2

- Count all pairs  $\{i, j\}$  that meet the conditions:
  1. Both  $i$  and  $j$  are frequent items.
  2. The pair  $\{i, j\}$ , hashes to a bucket number whose bit in the bit vector is 1.
- Notice all these conditions are necessary for the pair to have a chance of being frequent.

Data Science

82

82

## Memory Details

---

- Hash table requires buckets of 2-4 bytes.
  - Number of buckets thus almost 1/4-1/2 of the number of bytes of main memory.
- On second pass, a table of (item, item, count) triples is essential.
  - Thus, hash table must eliminate 2/3 of the candidate pairs to beat a-priori.

Data Science

83

83

## Multistage Algorithm

---


- **Key idea:** After Pass 1 of PCY, rehash only those pairs that qualify for Pass 2 of PCY.
- On middle pass, fewer pairs contribute to buckets, so fewer *false positives* --- frequent buckets with no frequent pair.

Data Science

84

84

## Multistage Picture



Data Science

85

85

## Multistage --- Pass 3

- Count only those pairs  $\{i, j\}$  that satisfy:
  1. Both  $i$  and  $j$  are frequent items.
  2. Using the first hash function, the pair hashes to a bucket whose bit in the first bit-vector is 1.
  3. Using the second hash function, the pair hashes to a bucket whose bit in the second bit-vector is 1.

Data Science

86

86

## Important Points

---

1. The two hash functions have to be independent.
2. We need to check both hashes on the third pass.
  - If not, we would wind up counting pairs of frequent items that hashed first to an infrequent bucket but happened to hash second to a frequent bucket.

Data Science

87

87

## Multihash

---

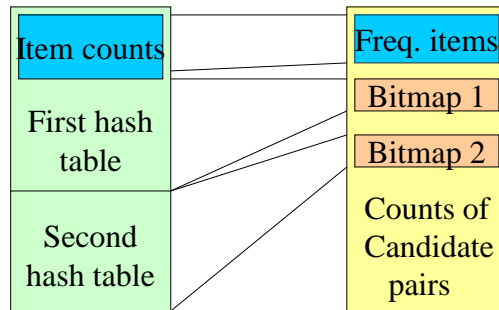
- **Key idea:** use several independent hash tables on the first pass.
- **Risk:** halving the number of buckets doubles the average count. We have to be sure most buckets will still not reach count  $s$ .
- If so, we can get a benefit like multistage, but in only 2 passes.

Data Science

88

88

## Multihash Picture



Data Science

89

89

## Extensions

- Either multistage or multihash can use more than two hash functions.
- In multistage, there is a point of diminishing returns, since the bit-vectors eventually consume all of main memory.
- For multihash, the bit-vectors total exactly what one PCY bitmap does, but too many hash functions makes all counts  $\geq s$ .

Data Science

90

90

## All (Or Most) Frequent Itemsets In $\leq 2$ Passes

---

- Simple algorithm.
- SON (Savasere, Omiecinski, and Navathe).
- Toivonen.

Data Science

91

91

## Simple Algorithm --- (1)

---


- Take a main-memory-sized random sample of the market baskets.
- Run a-priori or one of its improvements (for sets of all sizes, not just pairs) in main memory, so you don't pay for disk I/O each time you increase the size of itemsets.
  - Be sure you leave enough space for counts.

Data Science

92

92

## The Picture



Data Science

93

93

## Simple Algorithm --- (2)

- Use as your support threshold a suitable, scaled-back number.
  - E.g., if your sample is  $1/100$  of the baskets, use  $s/100$  as your support threshold instead of  $s$ .

Data Science

94

94

## Simple Algorithm --- Option

- Optionally, verify that your guesses are truly frequent in the entire data set by a second pass.
- But you don't catch sets frequent in the whole but not in the sample.
  - Smaller threshold, e.g.,  $s/125$ , helps.

Data Science

95

95

## Toivonen's Algorithm --- (1)

- Start as in the simple algorithm, but lower the threshold slightly for the sample.
  - **Example:** if the sample is 1% of the baskets, use  $s/125$  as the support threshold rather than  $s/100$ .
  - Goal is to avoid missing any itemset that is frequent in the full set of baskets.
- H. Toivonen. Sampling large databases for association rules. In *VLDB'96*

Data Science

96

96



## Toivonen's Algorithm --- (2)

---

- Add to the itemsets that are frequent in the sample the *negative border* of these itemsets.
- An itemset is in the negative border if it is not deemed frequent in the sample, but *all* its immediate subsets are.

Data Science

97

97

### Example: Negative Border

---

- $ABCD$  is in the negative border if and only if it is not frequent, but all of  $ABC$ ,  $BCD$ ,  $ACD$ , and  $ABD$  are.

Data Science

98

98

## Toivonen's Algorithm --- (3)

---

- In a second pass, count all candidate frequent itemsets from the first pass, and also count the negative border.
- If no itemset from the negative border turns out to be frequent, then the candidates found to be frequent in the whole data are *exactly* the frequent itemsets.

Data Science

99

99

## Toivonen's Algorithm --- (4)

---

- What if we find something in the negative border is actually frequent?
- We must start over again!
- Try to choose the support threshold so the probability of failure is low, while the number of itemsets checked on the second pass fits in main-memory.

Data Science

100

100

## Theorem:

---

- If there is an itemset frequent in the whole, but not frequent in the sample, then there is a member of the negative border frequent in the whole.

Data Science

101

101

## Proof:

---

- Suppose not; i.e., there is an itemset  $S$  frequent in the whole, but not frequent or in the negative border in the sample.
- Let  $T$  be a **smallest** subset of  $S$  that is not frequent in the sample.
- $T$  is frequent in the whole (monotonicity).
- $T$  is in the negative border (else not "smallest").

Data Science

102

102

## SON Algorithm --- (1)

---

- Repeatedly read small subsets of the baskets into main memory and perform the first pass of the simple algorithm on each subset.
- An itemset becomes a candidate if it is found to be frequent in *any* one or more subsets of the baskets.
- A. Savasere, E. Omiecinski, and S. Navathe. An efficient algorithm for mining association in large databases. In VLDB'95

Data Science

103

103

## SON Algorithm --- (2)

---

- On a second pass, count all the candidate itemsets and determine which are frequent in the entire set.
- Key “monotonicity” idea: an itemset cannot be frequent in the entire set of baskets unless it is frequent in at least one subset.

Data Science

104

104

## Compacting the Output

- A long pattern contains a combinatorial number of sub-patterns, e.g.,  $\{a_1, \dots, a_{100}\}$  contains ? sub-patterns

Data Science

105

105

## Compacting the Output

- A long pattern contains a combinatorial number of sub-patterns, e.g.,  $\{a_1, \dots, a_{100}\}$  contains  $\binom{100}{1} + \binom{100}{2} + \dots + \binom{100}{100} = 2^{100} - 1 = 1.27 \cdot 10^{30}$  sub-patterns!

Data Science

106

106

## Compacting the Output

- A long pattern contains a combinatorial number of sub-patterns, e.g.,  $\{a_1, \dots, a_{100}\}$  contains  $\binom{100}{1} + \binom{100}{2} + \dots + \binom{100}{100} = 2^{100} - 1 = 1.27 \cdot 10^{30}$  sub-patterns!
- Solution: Mine *closed patterns* and *max-patterns* instead
  1. *Maximal Frequent itemsets* : no immediate superset is frequent.
  2. *Closed itemsets* : no immediate superset has the same count.
    - Stores not only frequent information, but exact counts.

Data Science

107

107

## Closed Patterns and Max-Patterns

- An itemset  $X$  is *closed* if  $X$  is *frequent* and there exists *no super-pattern*  $Y \supset X$ , with the same support as  $X$  (proposed by Pasquier, et al. @ ICDT'99)
- An itemset  $X$  is a *max-pattern* if  $X$  is frequent and there exists no frequent super-pattern  $Y \supset X$  (proposed by Bayardo @ SIGMOD'98)
- Closed pattern is a lossless compression of freq. patterns
  - Reducing the # of patterns and rules

Data Science

108

108

## Example: Maximal/Closed

	Count	Maximal s=3	Closed
A	4	No	No
B	5	No	Yes
C	3	No	No
AB	4	Yes	Yes
AC	2	No	No
BC	3	Yes	Yes
ABC	2	No	Yes

Data Science

109

109

## Closed Patterns and Max-Patterns

- Exercise.  $DB = \{ \langle a_1, \dots, a_{100} \rangle, \langle a_1, \dots, a_{50} \rangle \}$ 
  - $Min\_sup = 1$ .
- What is the set of **closed itemsets**?

Data Science

110

110

## Closed Patterns and Max-Patterns

- Exercise.  $DB = \{ \langle a_1, \dots, a_{100} \rangle, \langle a_1, \dots, a_{50} \rangle \}$ 
  - $Min\_sup = 1.$
- What is the set of **closed itemsets**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$
  - $\langle a_1, \dots, a_{50} \rangle: 2$

Data Science

111

111

## Closed Patterns and Max-Patterns

- Exercise.  $DB = \{ \langle a_1, \dots, a_{100} \rangle, \langle a_1, \dots, a_{50} \rangle \}$ 
  - $Min\_sup = 1.$
- What is the set of **closed itemsets**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$
  - $\langle a_1, \dots, a_{50} \rangle: 2$
- What is the set of **max-patterns**?

Data Science

112

112



## Closed Patterns and Max-Patterns

- Exercise.  $DB = \{ \langle a_1, \dots, a_{100} \rangle, \langle a_1, \dots, a_{50} \rangle \}$ 
  - $Min\_sup = 1$ .
- What is the set of **closed itemsets**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$
  - $\langle a_1, \dots, a_{50} \rangle: 2$
- What is the set of **max-patterns**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$

Data Science

113

113

## Closed Patterns and Max-Patterns

- Exercise.  $DB = \{ \langle a_1, \dots, a_{100} \rangle, \langle a_1, \dots, a_{50} \rangle \}$ 
  - $Min\_sup = 1$ .
- What is the set of **closed itemsets**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$
  - $\langle a_1, \dots, a_{50} \rangle: 2$
- What is the set of **max-patterns**?
  - $\langle a_1, \dots, a_{100} \rangle: 1$
- What is the set of **all patterns**?

Data Science

114

114

## Ref: Basic Concepts of Frequent Pattern Mining

- (Association Rules) R. Agrawal, T. Imielinski, and A. Swami. Mining association rules between sets of items in large databases. SIGMOD'93.
- (Max-pattern) R. J. Bayardo. Efficiently mining long patterns from databases. SIGMOD'98.
- (Closed-pattern) N. Pasquier, Y. Bastide, R. Taouil, and L. Lakhal. Discovering frequent closed itemsets for association rules. ICDT'99.
- (Sequential pattern) R. Agrawal and R. Srikant. Mining sequential patterns. ICDE'95

Data Science

115

115

## Ref: Apriori and Its Improvements

- R. Agrawal and R. Srikant. Fast algorithms for mining association rules. VLDB'94.
- H. Mannila, H. Toivonen, and A. I. Verkamo. Efficient algorithms for discovering association rules. KDD'94.
- A. Savasere, E. Omiecinski, and S. Navathe. An efficient algorithm for mining association rules in large databases. VLDB'95.
- J. S. Park, M. S. Chen, and P. S. Yu. An effective hash-based algorithm for mining association rules. SIGMOD'95.
- H. Toivonen. Sampling large databases for association rules. VLDB'96.
- S. Brin, R. Motwani, J. D. Ullman, and S. Tsur. Dynamic itemset counting and implication rules for market basket analysis. SIGMOD'97.
- S. Sarawagi, S. Thomas, and R. Agrawal. Integrating association rule mining with relational database systems: Alternatives and implications. SIGMOD'98.

Data Science

116

116